

Uttkarsh Tyagi

+91 9193399551 | uttkarsh25tyagi@gmail.com | [linkedin.com/in/uttkarsh-tyagi](https://www.linkedin.com/in/uttkarsh-tyagi) | github.com/tyxgx | tyxgx.github.io/portfolio

Data & cloud-focused engineer building streaming data pipelines, cloud-native ML systems, and production backend services. Recently completed C-DAC's PG Certificate Programme in Big Data Analytics (PGCP-BDA). Comfortable across the stack: Python data/ML pipelines, AWS infrastructure via Terraform, and full-stack delivery.

EDUCATION

C-DAC — PG Certificate Programme in Big Data Analytics (PGCP-BDA)

Feb 2026 — Aug 2026

Uttarakhand Technical University, Dehradun

Sep 2021 — Jul 2025

B.Tech, Information Technology · CGPA 6.5

Delhi Public School, Ranipur, Haridwar

Class 12 — 88.4% · Class 10 — 86%

SKILLS

Languages Python, TypeScript/JavaScript, SQL

Data Engineering Apache Spark (Structured Streaming), Kafka/Redpanda, dbt, Delta Lake, medallion architecture (bronze/silver/gold), Apache Airflow

Cloud (AWS) Lambda, S3, API Gateway, Step Functions, EventBridge, Kinesis Firehose, Glue Data Catalog, ECR, IAM, X-Ray, Terraform (IaC)

Machine Learning scikit-learn, pandas, MLflow (tracking + registry), model evaluation & comparison pipelines, RAG pipelines (pgvector, SentenceTransformers), LLM API integration (Groq, Gemini, Ollama)

DevOps Docker, Docker Compose, GitHub Actions (CI/CD), Terraform, Prometheus, Grafana

Databases PostgreSQL, MongoDB, DynamoDB, schema design, indexing, query optimization

Tools Git/GitHub, Postman, Linux, REST APIs, WebSockets

PROJECTS

liveflights — Real-time flight intelligence platform

github.com/tyxgx/liveflights · [live demo](#)

Spark · Redpanda · Delta Lake · dbt · Terraform · AWS Lambda/S3/Step Functions · MLflow · FastAPI · Next.js

- Built an end-to-end streaming lakehouse — live flight data through Redpanda (Kafka-compatible) and Spark Structured Streaming into a medallion Delta Lake (bronze/silver/gold) on MinIO, modeled with dbt into Postgres marts.
- Designed and deployed a fully serverless replica of the same pipeline on AWS via Terraform: EventBridge-scheduled Lambda ingestion, Kinesis Firehose, S3 lake, Step Functions-orchestrated transforms, and a Glue Data Catalog — running live in production under a \$5/month budget.
- Root-caused AWS account-level provisioning restrictions (Glue Job/Crawler and CloudFront creation denied despite AdministratorAccess) using IAM policy simulation, and re-architected around them with zero downtime.
- Trained and served 4 purpose-built ML models (anomaly detection, forecasting, corridor/trajectory analysis) tracked via MLflow, exposed live through a WebSocket-backed FastAPI + Next.js ATC-style dashboard.

Global Music Intelligence Platform — Team Capstone (CDAC PGCP-BDA)

AWS Glue · Apache Spark · S3 · Athena · Terraform · GitHub Actions · PostgreSQL/pgvector · Django — 8-person team

- Independently built a parallel Bronze→Silver→Gold ETL pipeline in Terraform: Kaggle→EC2 ingestion, PySpark Glue jobs with incremental Silver loads, producing 5 Gold business tables — fully IaC, zero manual console steps, tested end-to-end.
- Debugged real infrastructure issues: Amazon Linux's `/tmp` being RAM-backed tmpfs (too small for a 10GB file), and `t3.small` burst-credit throttling under sustained CPU load (switched to non-burstable `m5.large`).
- Independently designed and built the platform's AI-powered natural-language layer — the RAG chatbot and Django web application (StreamPulse, below) — end to end.

StreamPulse — Spotify analytics dashboard + RAG chatbot

github.com/tyxgx/steampulse

Django · DRF · PostgreSQL · pgvector · SentenceTransformers · Groq/Gemini/Ollama · AWS EC2 · Nginx — independently built consumption layer for the platform above

- Built an ETL pipeline (boto3 + pyarrow + psycpg2) loading Hive-partitioned Parquet from an S3 Gold data lake into Postgres, and the full RAG pipeline from scratch: chunked Gold rows, embedded ~215K chunks locally (SentenceTransformers, 384-dim), indexed with pgvector's ivfflat.
- Built 3 retrieval paths — keyword-routed SQL, per-entity scoped retrieval, and vector search with a confidence gate that declines rather than hallucinates — plus a 3-tier LLM fallback (Groq → Gemini → Ollama).
- Iteratively debugged retrieval quality against 15 fixed test questions, flipping 7/15 from wrong to correct with 0 regressions. Deployed end-to-end on AWS EC2 (Postgres 17 + pgvector, Unicorn + Nginx, systemd).

Interactive ML — Stage-by-stage ML pipeline explorer

github.com/tyxgx/interactive-ml · [live demo](#)

FastAPI · scikit-learn · pandas · Next.js · TypeScript · Tailwind CSS

- Built an educational ML platform exposing every pipeline stage — load, split, preprocess, train, predict, evaluate — as an inspectable step, with leakage-safe preprocessing and a fair 5-algorithm comparison evaluated on an identical train/test split per session.

EXPERIENCE

Full Stack Developer Intern — Unlock Discounts, Bangalore (Remote)

Jul 2025 — Oct 2025

- Built a customer review section (React/JS/CSS) and APIs to fetch and display data from the company's database across platforms; created a certificate verification system for intern authentication.

Project Intern — BHEL, Haridwar

Jul 2025 — Aug 2025

- Developed a web portal (HTML/CSS/React, SQL backend) for equipment issuance with late tracking and recipient records.